

Neural Networks – Pre-Processing and Feature Extraction

Mohamed Krini

Christian-Albrechts-Universität zu Kiel
Faculty of Engineering
Institute of Electrical and Information Engineering
Digital Signal Processing and System Theory



Contents of the Lecture

Entire Semester

- ❑ Introduction
- ❑ *Pre-Processing and Feature Extraction*
- ❑ Threshold Logic Units - Single Perceptrons
- ❑ Multilayer Perceptrons
- ❑ Training Multilayer Perceptrons
- ❑ Radial Basis Function Networks
- ❑ Learning Vector Quantization
- ❑ Kohonen Self-Organizing Maps
- ❑ Hopfield and Recurrent Networks

Contents of this Part

Pre-Processing and Feature Extraction

- ❑ Introduction
- ❑ Input Normalization
- ❑ Feature Extraction for Speech Recognition
 - ❑ Analysis Filterbank
 - ❑ Mel Filterbank
 - ❑ Mel-filtered Cepstral Coefficients
 - ❑ Temporal Features
- ❑ Dimension Reduction

Pre-Processing and Feature Extraction


Introduction (1/3)

Classical Model of Pattern Recognition:

Task: Assign one of several classes to a set of measurements

Input: A vector of measurements (patterns)

Output: Number of the best class



Feature Extraction:

Extracts features from raw data (e.g. speech, image, weather data) that should ease the task of classification.

Classification:


The input features are assigned to one of K classes (e.g. speech signal is classified into phonemes and/or words).

Pre-Processing and Feature Extraction

Introduction (2/3)

Definitions and Properties:

- ❑ For most applications it is necessary first to *transform* the *data* into some new representation before presented to a network:



- ❑ The choice of *pre-processing* will be one of the most *significant factors* for determining the performance of the final system.
- ❑ In the simplest case the pre-processing performs a linear transformation of the input data (e.g. *input normalization*).

Pre-Processing and Feature Extraction

Introduction (3/3)

Definitions and Properties (continued):

- ❑ One of the most important forms of pre-processing involves a *reduction* in the *dimensionality* of the *input data*.
- ❑ In most situations a reduction in the dimensionality will result in *loss* of *information*.
 - ➡ *Goal*: Preserve as much of relevant information as possible.
- ❑ A network with fewer inputs has fewer adaptive parameters to be determined, leading to a network with better *generalization* properties.
- ❑ In addition a network with fewer weights may be *faster* to train.
- ❑ The transformed inputs are often called *features* and the process of generating them is called *feature extraction*.

Pre-Processing and Feature Extraction

Input Normalization

Input Normalization

Pre-Processing and Feature Extraction

Input Normalization (1/3)

Properties:

- ❑ Simple linear rescaling of the input variables is one of the *most common forms of pre-processing*.
- ❑ Often useful if different variables have typical values which differ significantly.
- ❑ Input normalization *avoids scaling problems*.
- ❑ The rescaling is useful for radial basis function networks and multilayer networks:
 - ➡ If variation in one parameter is small with respect to the others it will contribute very little to distance measures.
- ❑ Each of the input variables is treated independently for linear rescaling.

Pre-Processing and Feature Extraction

Input Normalization (2/3)

Procedure for Input Normalization:

- We calculate the *mean* and the *variance* with respect to the training set:

Training set consists of N -dimensional input vectors:

$$\mathbf{x}^{(m)} = \left[x_0^{(m)}, x_1^{(m)}, \dots, x_{N-1}^{(m)} \right]^T, \quad \text{with } m \in \{0, \dots, M-1\}.$$

← Number of training patterns

Mean:

$$\mu_i = \frac{1}{M} \sum_{m=0}^{M-1} x_i^{(m)}.$$

Variance:

$$\sigma_i^2 = \frac{1}{M} \sum_{m=0}^{M-1} \left(x_i^{(m)} - \mu_i \right)^2.$$


- Normalize the input vectors to *expected value 0* and *standard deviation 1*:

$$\tilde{x}_i^{(m)} = \frac{x_i^{(m)} - \mu_i}{\sigma_i}.$$

Pre-Processing and Feature Extraction

Input Normalization (3/3)

Experiment:



- Large data set consists of two dimensional feature vectors.
- Blue contour line represents feature vectors without normalization.
- Red contour line shows the input vectors after normalization.
- Mean value: X

Feature Extraction:

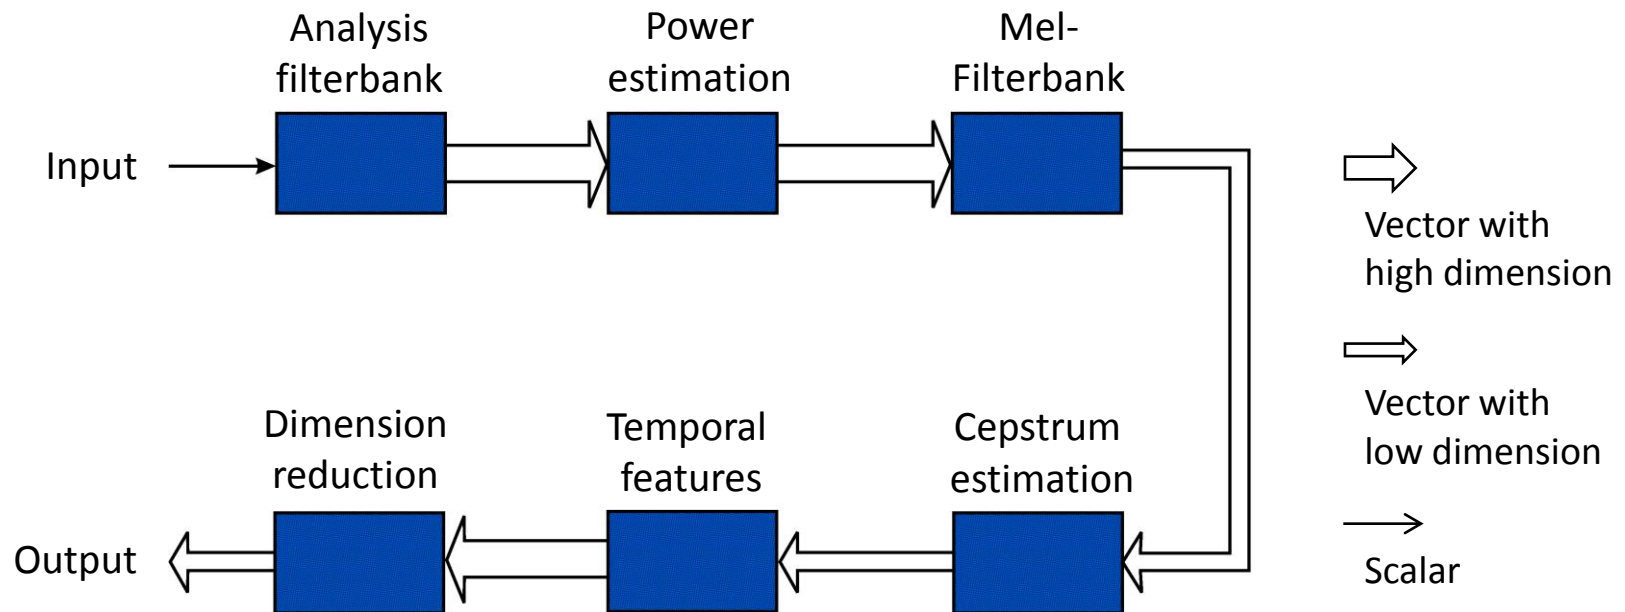
- Application for Speech Recognition***

Pre-Processing and Feature Extraction

Feature Extraction for Speech Recognition

Overview:

- As an application example we consider the pre-processing and feature extraction for *speech recognition*.



Pre-Processing and Feature Extraction

Feature Extraction – Analysis Filterbank (1/3)

Analysis Filterbank:

- Input signal is first **segmented** into blocks of appropriate size:

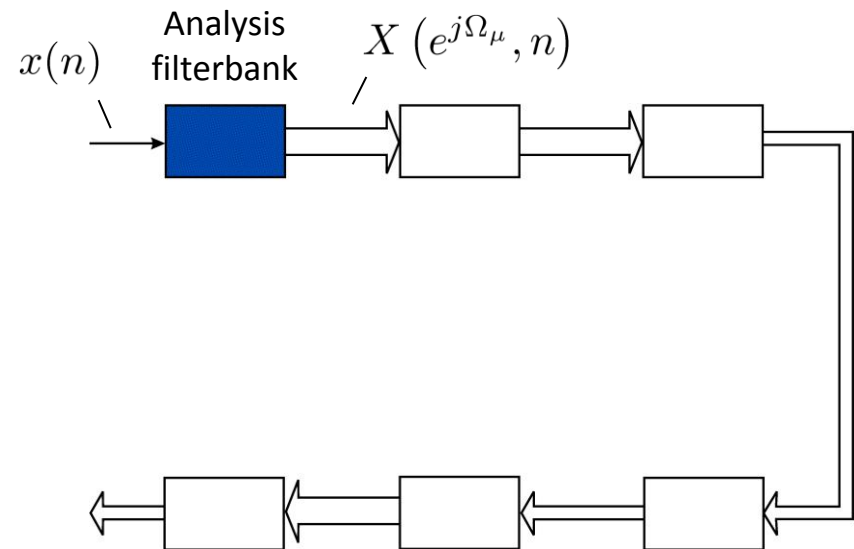
$$\tilde{x}(n) = [x(n), x(n-1), \dots, x(n-N+1)]^T.$$

- Usually adjacent input segments are **overlapped** (modeled by a **subsampling** factor R):

$$x(nR) = \tilde{x}(nR).$$

- By applying a window function h_k and computing the DFT, the **short-term spectrum** results:

$$X(e^{j\Omega_\mu}, n) = \sum_{k=0}^{N-1} x(nR-k) h_k e^{-j\Omega_\mu k}, \quad \text{with } \mu \in \{0, \dots, N-1\}.$$



Pre-Processing and Feature Extraction

Feature Extraction – Analysis Filterbank (2/3)

Principle:

Used parameters:

Block length:

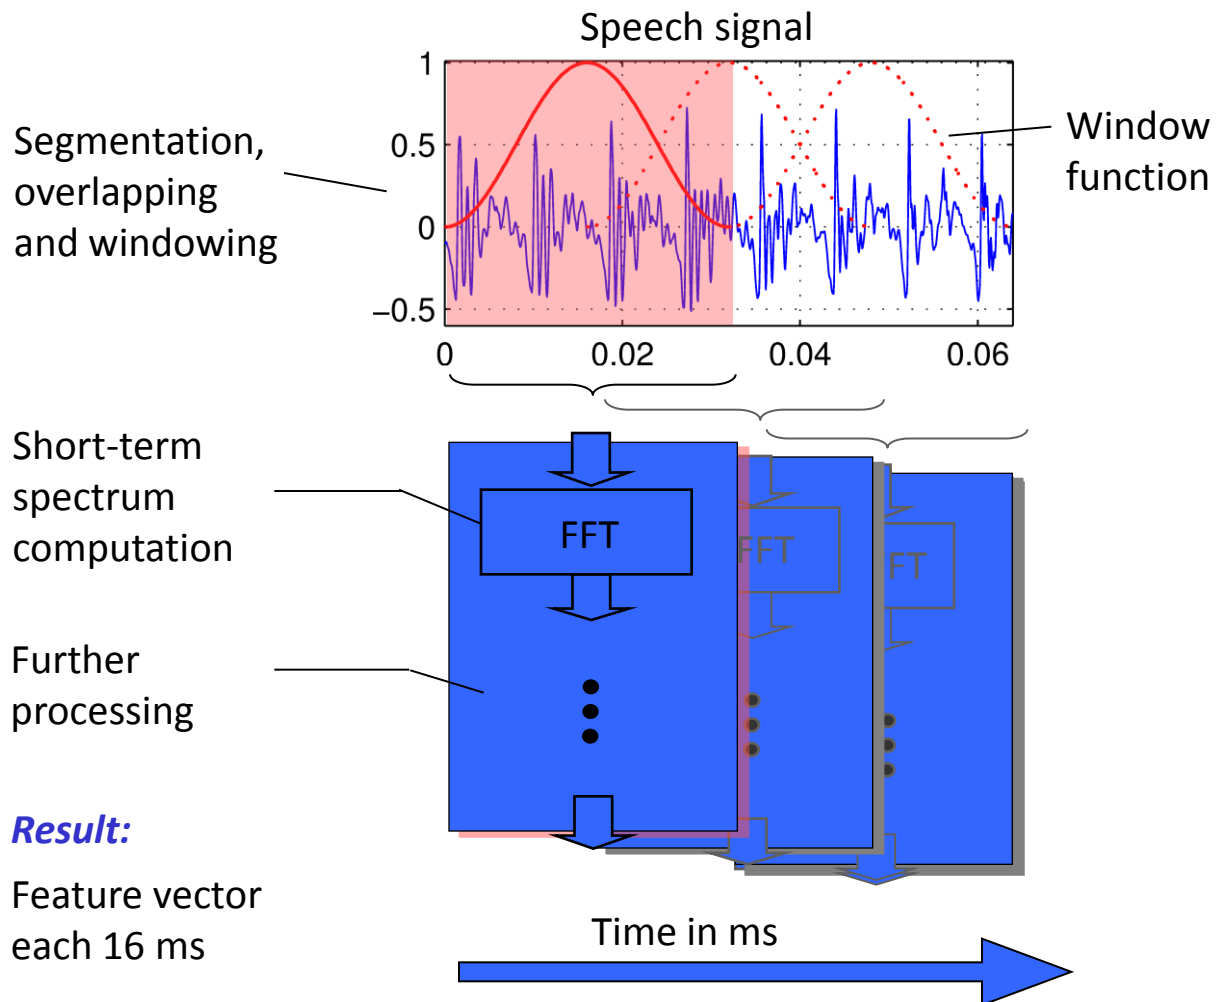
$$N = 512$$

Subsampling rate:

$$R = 256$$

Sampling rate:

$$f_s = 16 \text{ kHz}$$





Result:

Feature vector
each 16 ms

Pre-Processing and Feature Extraction

Feature Extraction – Analysis Filterbank (3/3)

Example of Time-Frequency Representation:



Pre-Processing and Feature Extraction

Feature Extraction – Power Estimation

Power Estimation:

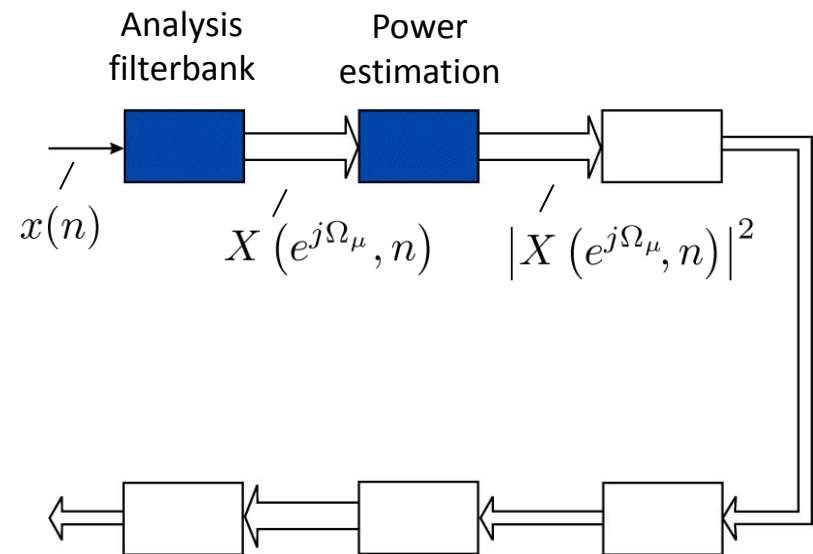
- Absolute square of the input spectrum:

$$|X(e^{j\Omega_\mu}, n)|^2 = \Re^2 \{X(e^{j\Omega_\mu}, n)\} + \Im^2 \{X(e^{j\Omega_\mu}, n)\}.$$

- **Approximation** of the absolute value (low computational cost needed, low dynamic):

$$|X(e^{j\Omega_\mu}, n)| \approx K \left| \Re \{X(e^{j\Omega_\mu}, n)\} \right| + K \left| \Im \{X(e^{j\Omega_\mu}, n)\} \right|.$$

- **Amplitude** is much **more important than** the **phase** (phase is discarded), the results are real numbers.



Pre-Processing and Feature Extraction

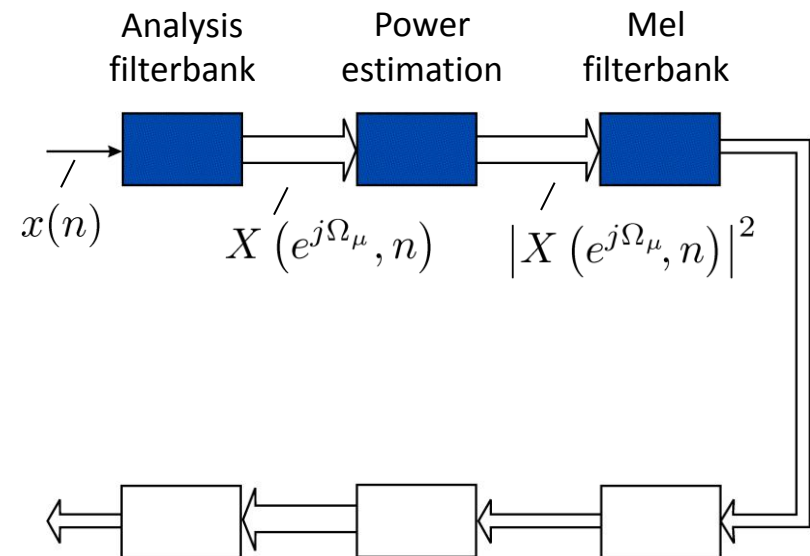
Feature Extraction – Mel Filterbank (1/5)

Mel Filterbank:

- A set of *triangular filters* is used to approximate the frequency resolution of the human ear.
- Approximated formula to compute the frequency (or pitch) in Mel for a given frequency f in Hz:

$$m = 2595 \text{ Mel} \log_{10} \left(1 + \frac{f}{700 \text{ Hz}} \right) .$$

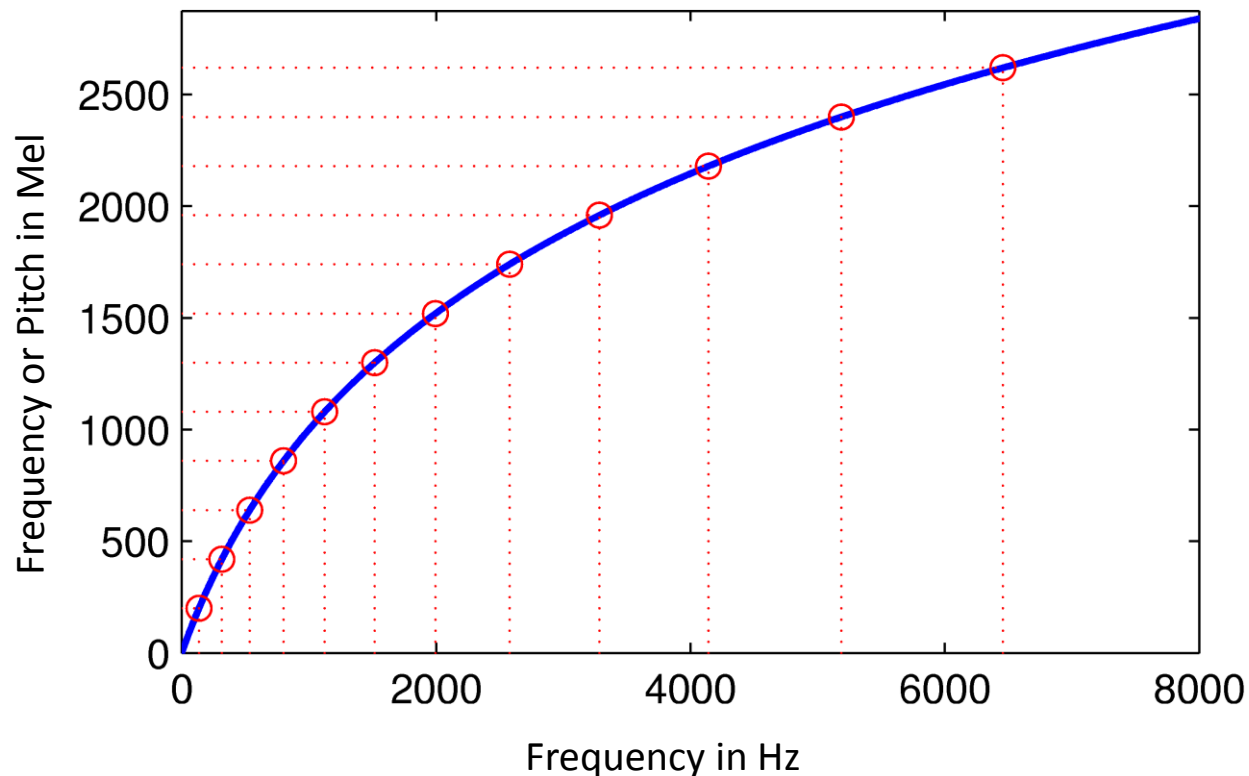
- The Mel-frequency scale is linear up to 1 kHz and logarithmic thereafter.



Pre-Processing and Feature Extraction

Feature Extraction – Mel Filterbank (2/5)

Mel-Scale Characteristic:

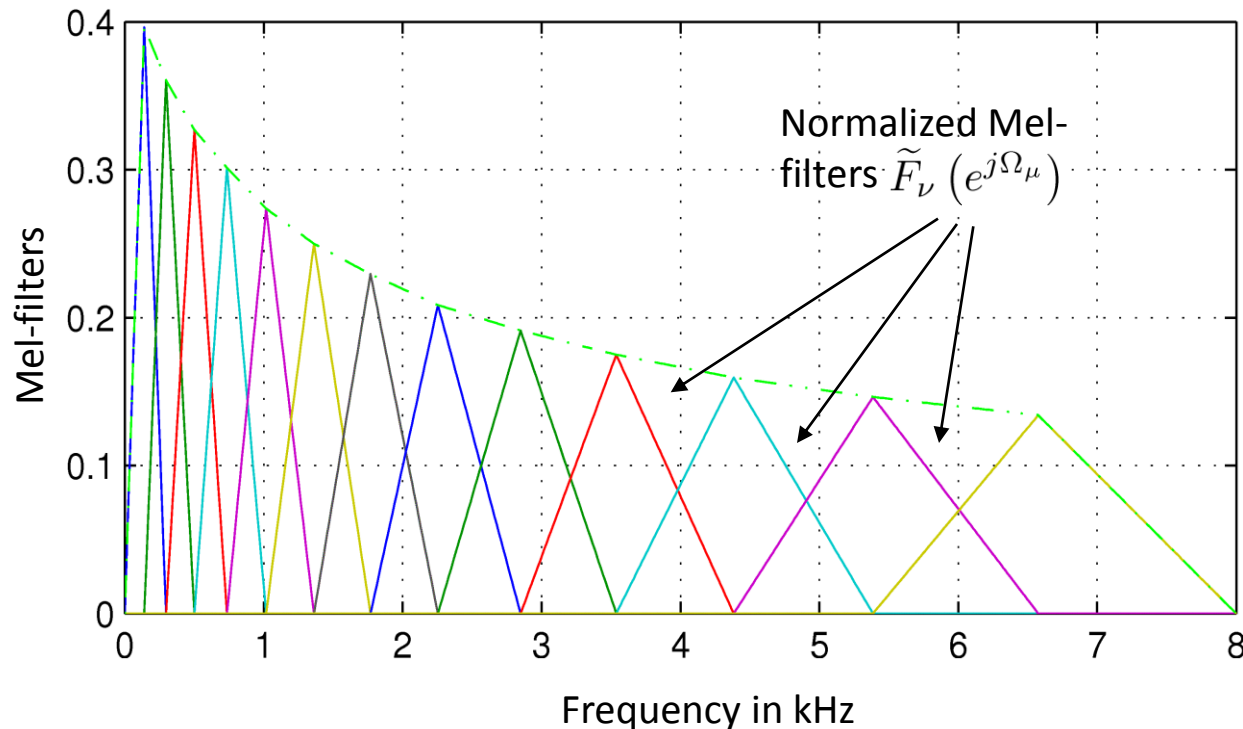


- The Mel-frequency scale shows a **logarithmic** behavior.
- Mel-frequency axis (Mel scale) divided into 13 **equispaced** bands in this example.
- A set of **overlapping Mel filters** is designed such that the center frequencies of the filters are equidistant on the Mel scale.

Pre-Processing and Feature Extraction

Feature Extraction – Mel Filterbank (3/5)

Overlapping Mel-Filters:



- A set of **triangular** filter banks is used to approximate the frequency resolution of the human ear.
- Typically 15 up to 30 Mel filters are applied for a sampling rate range from 8 to 16 kHz.
- The power of the spectrum is mapped onto the Mel scale using triangular **overlapping** windows.

Pre-Processing and Feature Extraction

Feature Extraction – Mel Filterbank (4/5)

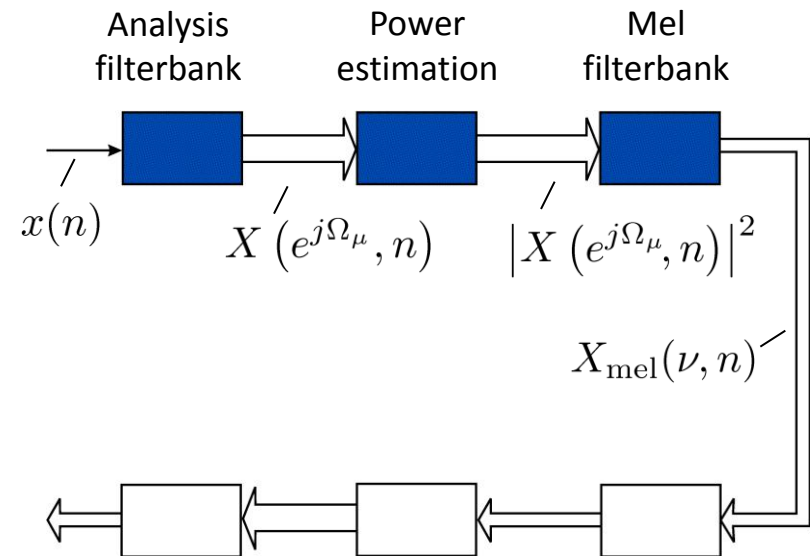
Power Estimation in Mel Domain:

- The **Mel power spectrum** is determined by:

$$X_{\text{mel}}(\nu, n) = \frac{\sum_{\mu=0}^{N-1} F_{\nu}(e^{j\Omega_{\mu}}) |X(e^{j\Omega_{\mu}}, n)|^2}{\sum_{\mu=0}^{N-1} F_{\nu}(e^{j\Omega_{\mu}})},$$

with $\nu \in \{0, \dots, M-1\}$.


- Mel filters are **normalized** to guarantee the same power at the output in case of white noise excitation.
- The outputs are **real-valued** and **reduced in dimension** compared to the input.



Pre-Processing and Feature Extraction

Feature Extraction – Mel Filterbank (5/5)

Example:



- The dimension of the input features is reduced after transformation into the Mel-domain.

Pre-Processing and Feature Extraction

Feature Extraction – Cepstrum (1/3)

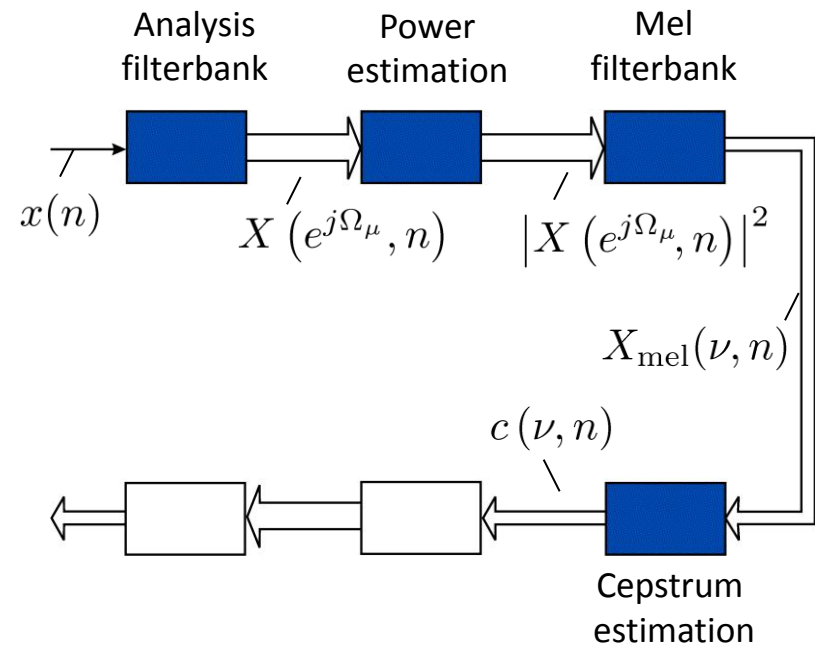
Cepstrum:

- The **cepstrum** is defined as the IDFT of the logarithm of the power spectrum:

$$c(\nu, n) = \frac{1}{M} \sum_{k=0}^{M-1} \log \left\{ X_{\text{mel}}(k, n) \right\} e^{j \frac{2\pi}{M} k \nu},$$

$$\mathbf{c}(n) = [c(0, n), \dots, c(M-1, n)]^T.$$

- Useful transformation for decorrelating and removing speaker dependent information from the input features.
- The resulting cepstrum is symmetric, it's therefore sufficient to use only the first half of the cepstral coefficients.



Pre-Processing and Feature Extraction

Feature Extraction – Cepstrum (2/3)

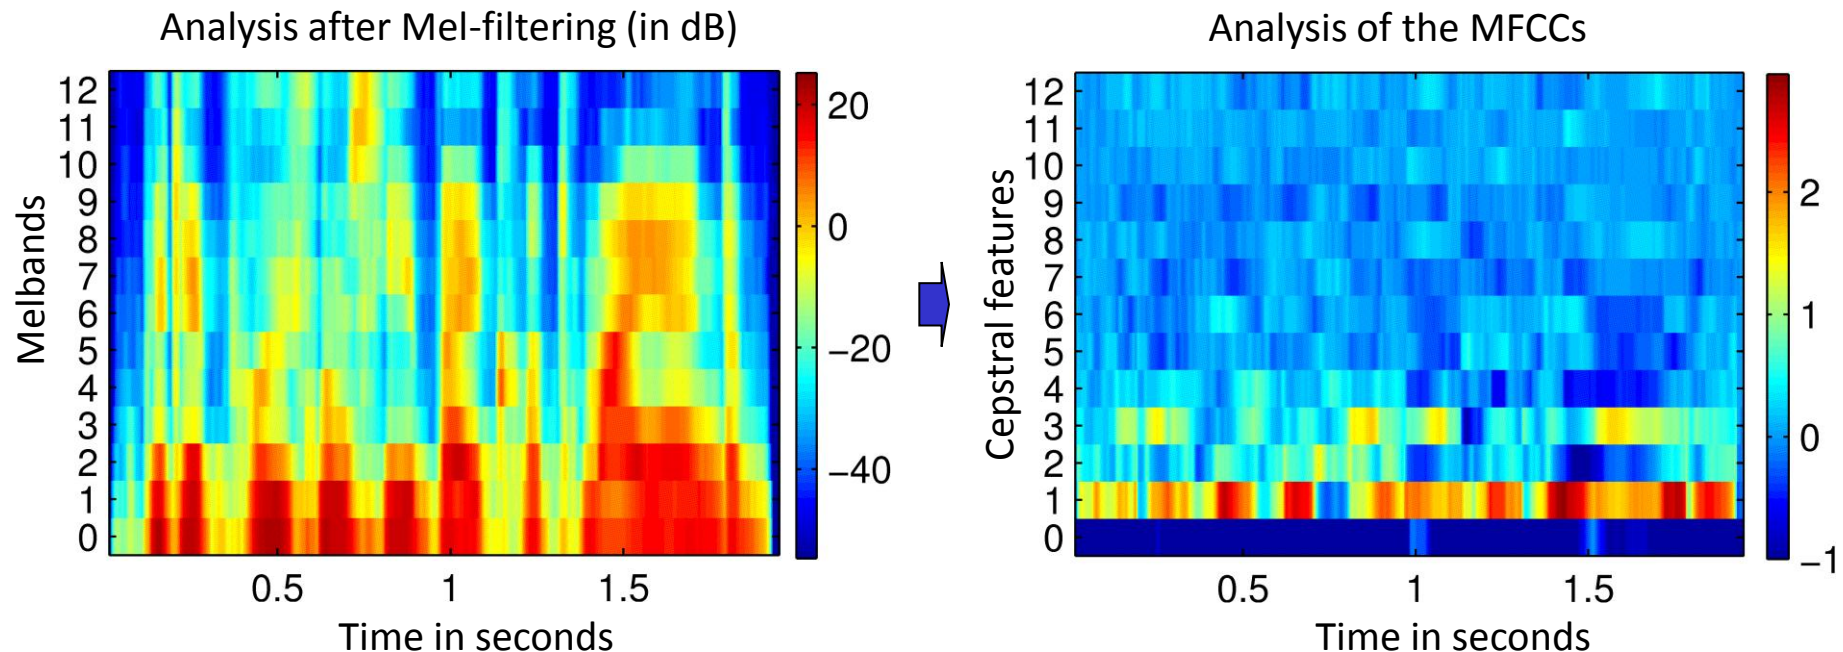
Properties:

- ❑ The purpose of the transformation is to *decorrelate* the logarithmic input features:
 - ➡ Good for classification.
- ❑ The outcome features are called *Mel-Frequency Cepstral Coefficients* (MFCCs).
- ❑ Since the input vector is real-valued the IDFT can be replaced by IDCT for efficient implementation.
- ❑ The cepstrum contains largely *vocal tract* information (concentrates at lower bands).
- ❑ Usually the *dimension* of the *feature vector* is *reduced* by minimizing the behavior of the *pitch frequency* (concentrates at higher bands). Often the last third of the feature elements is discarded.
- ❑ Cepstral coefficients are *intensive to loudness* (only energy changes, cepstrum unaffected).

Pre-Processing and Feature Extraction

Feature Extraction – Cepstrum (3/3)

Example:



- The computed *MFCCs* of this example have mainly high energy at lower quefrencies.

Pre-Processing and Feature Extraction

Feature Extraction – Temporal Features (1/2)

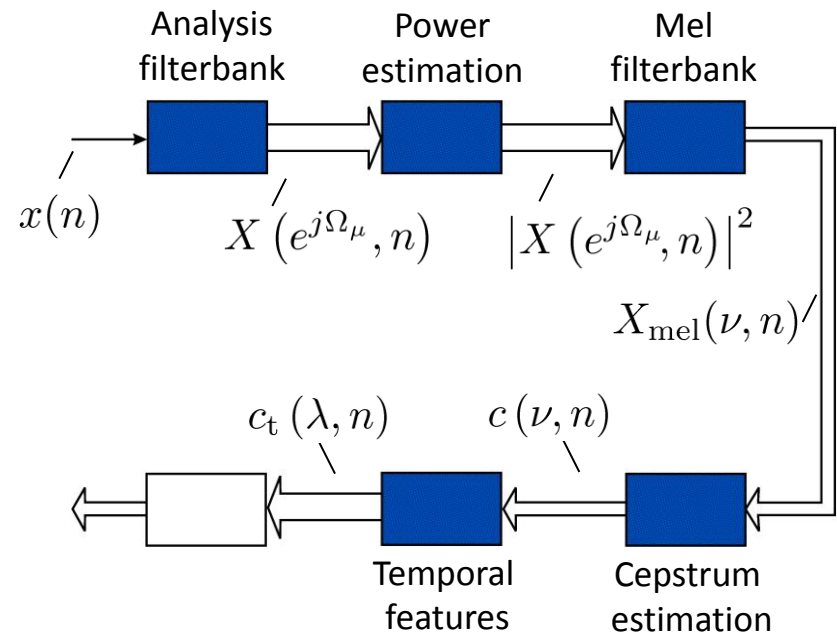
Temporal Features:

- After feature extraction often some number of successive **feature vectors** are **combined**.
- In some cases the difference of adjacent feature vectors is computed (called **delta features**) or the difference of two adjacent differences (called **delta-delta features**).
- In the first case, successive vectors are **stacked** to a **multi-feature vector** as follows:

$$\mathbf{c}_t(n) = [\mathbf{c}^T(n - P_b), \dots, \mathbf{c}^T(n), \dots, \mathbf{c}^T(n + P_f)]^T,$$

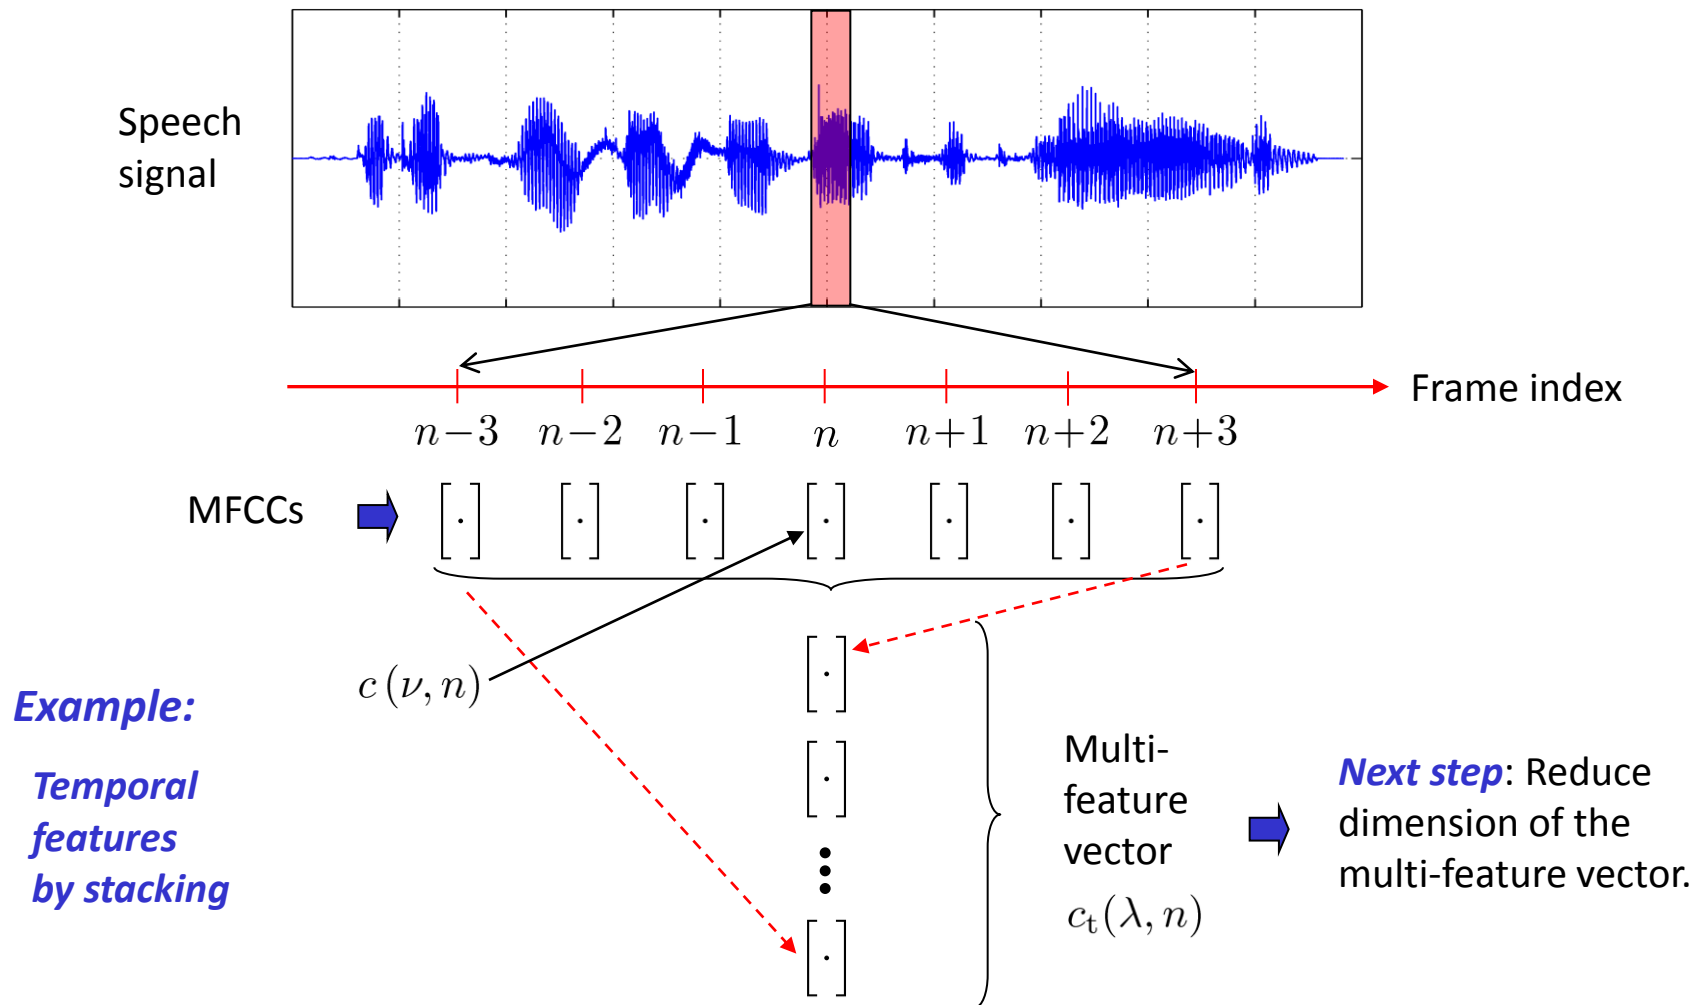
Number of feature vectors from the past and the future

$$\mathbf{c}_t(n) = [\mathbf{c}_t(0, n), \dots, \mathbf{c}_t(M_t - 1, n)]^T, \quad M_t = M + M(P_f + P_b).$$



Pre-Processing and Feature Extraction

Feature Extraction – Temporal Features (2/2)

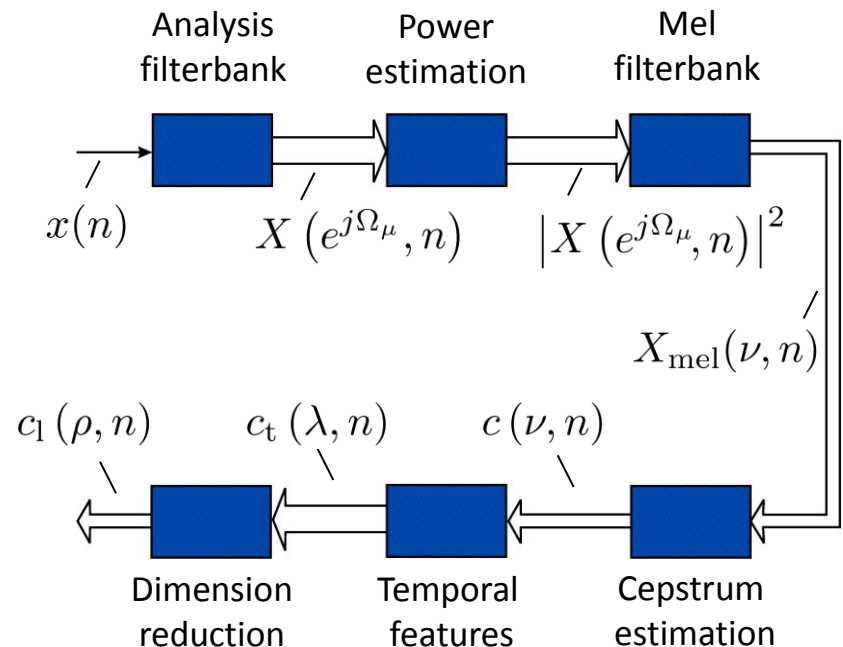


Pre-Processing and Feature Extraction

Feature Extraction – Feature Space Transformation (1/6)

Dimension Reduction:

- ❑ A **feature space transformation** is performed to reduce the dimension of the input features $c_t(n)$.
- ❑ The so-called **linear discriminant analysis** (LDA) can be applied.
- ❑ The objective of LDA is to perform dimensionality reduction while **preserving** as much of the **class discriminatory** information as possible.
- ❑ The **variance** of features which corresponds to a feature class is **minimized** while the **distance** between feature **classes** is **maximized**.

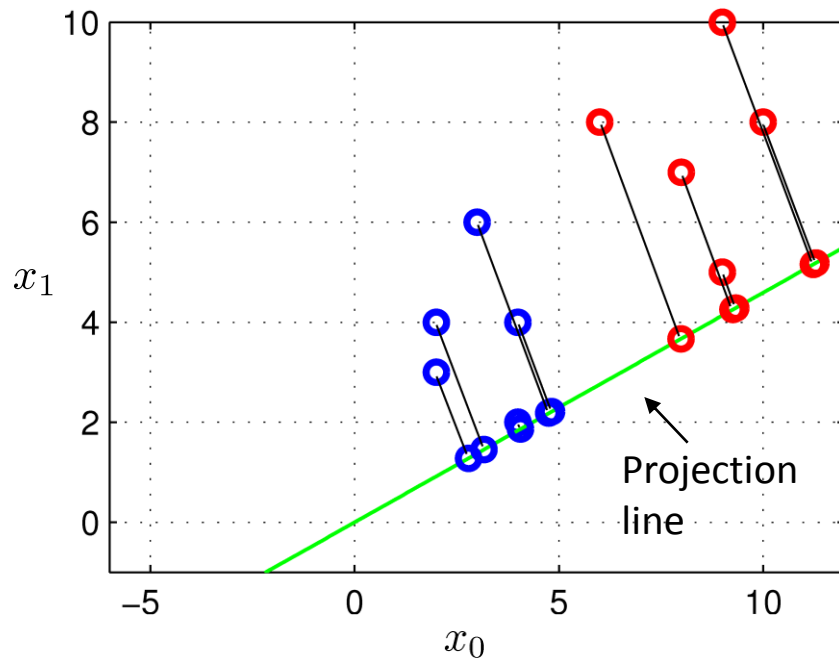


Pre-Processing and Feature Extraction


Feature Extraction – Feature Space Transformation (2/6)

LDA Example:

Good separability between two classes



Bad separability between two classes



□ Select the line that *maximizes* the *separability* of the *projected features*.

Pre-Processing and Feature Extraction

Feature Extraction – Feature Space Transformation (3/6)

LDA Derivation – Two Classes:

- Suppose we have a data set T consisting of M observations of a N - dimensional Euclidian variable \mathbf{x} :

$$T = \left\{ \mathbf{x}^{(0)}, \mathbf{x}^{(1)}, \dots, \mathbf{x}^{(M-1)} \right\}, \quad \text{with} \quad \mathbf{x}^{(m)} = \left[x_0^{(m)}, x_1^{(m)}, \dots, x_{N-1}^{(m)} \right]^T.$$

- We take the N -dimensional input vector \mathbf{x} and **project** it down to one dimension using a projection vector \mathbf{w} :

$$y^{(m)} = \mathbf{w}^T \mathbf{x}^{(m)}, \quad \text{with} \quad \mathbf{w} = [w_0, w_1, \dots, w_{N-1}]^T.$$

- In general the projection onto one dimension leads to a considerable loss of information:

➡ ***Select a projection that maximizes the class separation.***

Pre-Processing and Feature Extraction

Feature Extraction – Feature Space Transformation (4/6)

LDA Derivation (continued):

- The simplest measure of the separation of the classes is the separation of the **projected class means**. For a **two class problem**, chose w so as to maximize:

$$\tilde{\mu}_1 - \tilde{\mu}_0 = w^T (\mu_1 - \mu_0) \rightarrow \max,$$

$$\text{with } \mu_0 = \frac{1}{M_0} \sum_{l \in C_0} x^{(l)}, \quad \mu_1 = \frac{1}{M_1} \sum_{l \in C_1} x^{(l)}.$$

Mean vectors of the two classes C_0 and C_1 with M_0 and M_1 points.

- The within **class variances** of the transformed data from classes C_0 and C_1 are given by:

$$s_0^2 = \sum_{l \in C_0} (y^{(l)} - \tilde{\mu}_0)^2, \quad s_1^2 = \sum_{l \in C_1} (y^{(l)} - \tilde{\mu}_1)^2.$$

Transformed
feature

- The **total** within **class variance** is simply: $s_0^2 + s_1^2$.

Pre-Processing and Feature Extraction

Feature Extraction – Feature Space Transformation (5/6)

LDA Derivation (continued):

- ❑ **Fisher's idea:** Maximize a function that will give a large separation between the projected class **means** while also giving a small **variance** within each class, thereby minimizing the class overlap. The Fisher criterion is defined as:

$$J(\mathbf{w}) = \frac{(\tilde{\mu}_1 - \tilde{\mu}_0)^2}{s_0^2 + s_1^2} = \frac{\mathbf{w}^T \mathbf{S}_b \mathbf{w}}{\mathbf{w}^T \mathbf{S}_w \mathbf{w}}.$$

- ❑ The matrix \mathbf{S}_b is called **between class covariance matrix**:

$$\mathbf{S}_b = (\boldsymbol{\mu}_1 - \boldsymbol{\mu}_0) (\boldsymbol{\mu}_1 - \boldsymbol{\mu}_0)^T.$$

- ❑ The matrix \mathbf{S}_w is called **within class covariance matrix**:

$$\mathbf{S}_w = \sum_{l \in C_0} (\mathbf{x}^{(l)} - \boldsymbol{\mu}_0) (\mathbf{x}^{(l)} - \boldsymbol{\mu}_0)^T + \sum_{l \in C_1} (\mathbf{x}^{(l)} - \boldsymbol{\mu}_1) (\mathbf{x}^{(l)} - \boldsymbol{\mu}_1)^T.$$

Pre-Processing and Feature Extraction

Feature Extraction – Feature Space Transformation (6/6)

LDA Derivation (continued):

- The Fishers *criterion* is *maximized* to find the optimal weights:

$$\mathbf{w}_{\text{opt}} = \underset{\mathbf{w}}{\operatorname{argmax}} \left(\frac{\mathbf{w}^T \mathbf{S}_b \mathbf{w}}{\mathbf{w}^T \mathbf{S}_w \mathbf{w}} \right) .$$

- Differentiating $J(\mathbf{w})$ with respect to \mathbf{w} and setting the equation to zero we obtain:

$$\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}} = \frac{\partial}{\partial \mathbf{w}} \left(\frac{\mathbf{w}^T \mathbf{S}_b \mathbf{w}}{\mathbf{w}^T \mathbf{S}_w \mathbf{w}} \right) = \mathbf{0} .$$

- The result is known as *Fishers linear discriminant*:

$$\mathbf{w}_{\text{opt}} = \underset{\mathbf{w}}{\operatorname{argmax}} \left(\frac{\mathbf{w}^T \mathbf{S}_b \mathbf{w}}{\mathbf{w}^T \mathbf{S}_w \mathbf{w}} \right) = \mathbf{S}_w^{-1} (\boldsymbol{\mu}_1 - \boldsymbol{\mu}_0) .$$

Pre-Processing and Feature Extraction

Literature

Further details can be found in:

- ❑ C. Bishop: *Pattern Recognition and Machine Learning*, Springer, Berlin, Germany, 2006.
- ❑ C. Bishop: *Neural Networks for Pattern Recognition*, Oxford University Press, UK, 1996.
- ❑ E Schukat-Talamanzzini: *Automatische Spracherkennung – Grundlagen, Statistische Modelle und effiziente Algorithmen*, Vieweg, 1995.
- ❑ L. Rabiner, B.-H. Juang: *Fundamentals of Speech Recognition*, Prentice-Hall, 1993.